

Survival Analysis of Heart Failure Patients

Souradeep Sen

Department of Computer Science,

University of Exeter

(Dated: July, 2023)

The study aims to compare the performance of Deep Learning (DL) architectures for predicting mortality and hospitalization in heart failure (HF) patients, against traditional survival analysis techniques. The aim is to see if combining unsupervised and supervised learning can help estimate survival probabilities based on contextual and historic features from clinical data, and how these predictions fare against traditional techniques. By leveraging longitudinal patient data made available in the form of Electronic Health Records (EHR), a new perspective is sought on the performance of machine learning risk prediction models compared to conventional survival analysis in HF patients.

1. INTRODUCTION

Heart failure is a clinical syndrome interfering with the heart's ability to pump blood, leading to a reduced ability to perform systemic circulation. This condition is widespread globally - approximately 26 million people worldwide are estimated to be affected by heart failure [1]. Hence, accurately predicting risk is crucial for improving patient outcomes. Traditional survival analysis can be limited in its ability to handle complex non-linear dependencies as well as accounting for time-variant patient characteristics [this is not really true]. Deep Learning is an exciting tool in this aspect due to its innate ability to handle complex non-linear relationships [include citation - NNs are universal function approximators]. While there has been work done on adapting deep learning to the survival analysis domain (see:), it is still a growing field. As DL is generally labeled as data-hungry [include citation], the advent of EHR data with its vast longitudinal bandwidth(?) looks promising to be used in conjunction with DL. This study will attempt to validate the hypothesis that HDL outperforms traditional survival analysis in terms of prediction accuracy using real-world EHR data. The findings may have important implications for clinical practice, healthcare resource allocation, and future research in risk prediction modeling for HF patients with frailty.

The paper is organized as follows. Section 2 contains a brief review of related work in this capacity. Section 3 presents a more rigorous understanding of survival analysis, before building up to how discrete hazard rates [2] [CHECK THIS?] (as used in this paper) can be parameterized by a neural network [3]. A suitable loss function is chosen from the available literature and is derived as per [3]. Section 4 delves into Uncertainty Quantification, by way of Monte Carlo (MC) dropout and looks at explainability of the model(s) built. Section 5 discusses the data used for the paper - MIMIC-IV. Section 6 looks at the results by means of experiments over real and synthetic data and compares the architecture to existing solutions from deep learning and traditional survival analysis. Section 7 is reserved for proposed extensions to the model and further work that could be done. Some of the su-

pervised and unsupervised methods used in this paper (MLP, CNN, PCA) are detailed in the Appendices.

2. RELATED WORK

Traditional survival analysis has been used extensively to predict mortality in this patient population. Deep learning methods have also been employed - see [4], [5], [6], [7] and [8]. However, limited work has been done in predicting mortality and hospitalization in HF patients with frailty, especially using electronic health records (EHR) data. Several past papers have addressed predictive modeling for heart failure patients. A deep neural network model with learned medical feature embedding is proposed in [9] to address high dimensionality and temporality in electronic health record (EHR) data. Here, a convolutional neural network is used to capture non-linear longitudinal evolution of EHRs and local temporal dependency for risk prediction, and embed medical features to account for high dimensionality. Experiments show promising results in predicting risks for congestive heart failure.

Personalized predictive modeling is investigated in [10], which aims to build specific models for individual patients using similar patient cohorts to capture their specific characteristics. According to this study, although CNNs have shown promise on measuring patient similarity, one disadvantage is that they could not utilize temporal and contextual information of EHRs. To measure patient similarity using EHRs, the authors proposed a time-fusion CNN framework. A vector representation was generated for each patient, which was then utilized for measuring patient similarity and personalized disease prediction. Dynamic updates to a CNN model are explored in [11] as more data is gathered over time - this architecture lends itself well to real-time mortality risk prediction.

Maintaining interpretability across deep learning models is explored in [12]. Many previous studies using machine learning for modeling the risk of HF in patients have fo-

cused on discretized outputs. This study aims to consider incidences as time-to-event to enable continuous probabilistic risk prediction for hospitalization and mortality, addressing a critical need in patient care. The use of EHR such as those available in CPRD, allows access to comprehensive longitudinal data, which captures the entire cycle of a patient's diagnosis and treatment. Add citations for PyCox, Deep Survival Machines, Deep Survival Analysis, Faraggi-Simon, DeepSurv, RNN-SURV, etc.

3. SURVIVAL ANALYSIS

A. Basics

Survival analysis deals with the estimation of a survival distribution representing the probability of an event of interest, typically a failure, to occur beyond a certain time in the future [13]. One way to specify the survival distribution is through the survival function. The survival function defines the probability of surviving till point t [14].

$$S(t) = P(T > t), \quad 0 < t < \infty$$

It can be thought of as the complement of the cumulative distribution function $F(t)$.

$$S(t) = P(T > t) = 1 - P(T \leq t) = 1 - F(t), \quad 0 < t < \infty$$

Another way to specify the survival distribution is through the hazard function, which denotes the instantaneous rate of failure.

$$h(t) = \lim_{\delta \rightarrow 0} \frac{P(t < T < t + \delta | T > t)}{\delta}$$

For the continuous-time scenario, the hazard function and survival function are related as follows.

$$\begin{aligned} f(t) &= \frac{d}{dt}F(t) = -\frac{d}{dt}S(t) \\ h(t) &= \frac{f(t)}{S(t)} \end{aligned}$$

where $f(t)$ is the probability mass function (PMF). This says that the hazard is the probability of the subject experiencing an event at time t , provided that the subject is alive till time t . It can be further simplified as

$$\begin{aligned} h(t) &= -\frac{d}{dt}S(t) \frac{1}{S(t)} \\ \implies S(t) &= \exp\left(-\int_0^t h(u)du\right) \end{aligned}$$

This relationship produces the survival function from a hazard function [14].

Before moving to the discrete setting, some formal notation for the data is introduced. The data is assumed to be right-censored. Hence, the data, \mathcal{D} can be represented as a set of tuples $\{(x_i, t_i, d_i)\}_{i=1}^N$ [13]. Here, $x_i \in \mathbb{R}^d$ are covariates for patient i . t_i is the time of an event or censoring such that $t_i = \min(T_i, C_i)$, where T_i and C_i respectively denote the times of event and censoring. A subject is assumed to have either experienced an event or have been censored, but not both. d_i is an indicator that signifies whether t_i is event time or censoring time. $d = 1$ for a subject that experiences the event (uncensored) while $d = 0$ for a subject that is censored before experiencing the event. More formally, $d = \mathbb{I}\{T_i \leq C_i\}$. Later in the paper, experiments with time-variant covariates will necessitate the use of a null masking matrix, \mathcal{M} [cite DeepHit].

B. Discrete-Time Survival Analysis

For hazard and survival calculation in a discrete-time setting, the following formulation from [3] and earlier [2] is presented. Let $\mathbb{T} = \{\tau_1, \tau_2, \dots\}$ denote the timestamps, i.e. the indices of the discrete times corresponding to different subjects in the data. The event time is $T^* \in \mathbb{T}$. The definitions of PMF and survival function follow as

$$\begin{aligned} f(\tau_j) &= P(T^* = \tau_j), \\ S(\tau_j) &= P(T^* > \tau_j) = \sum_{k>j} f(\tau_k) \end{aligned}$$

It can be seen that the hazard at time τ_j $h(\tau_j)$, is just the probability of an event happening at time τ_j , given the subject has survived till the previous time step τ_{j-1} .

$$h(\tau_j) = P(T^* = \tau_j | T^* > \tau_{j-1}) = \frac{f(\tau_j)}{S(\tau_{j-1})} \quad (1)$$

$$\implies h(\tau_j) = \frac{S(\tau_{j-1}) - S(\tau_j)}{S(\tau_{j-1})} \quad (2)$$

$$\implies S(\tau_j) = (1 - h(\tau_j))S(\tau_{j-1}) \quad (3)$$

Recursively, the survival function can be parameterized wholly in terms of the hazard function as

$$S(\tau_j) = \prod_{k=1}^j (1 - h(\tau_k))$$

If there were no censoring involved, the likelihood of observing n failures (events) is of the form

$$L(t_1, t_2, \dots, t_n) = f(t_1)f(t_2) \dots f(t_n) = \prod_{i=1}^n f(t_i)$$

As per [14], for an observed event, the pdf is retained. But for a right-censored observation, it is replaced by

the survival function, as that observation is known only to exceed a particular value. The likelihood then becomes

$$L(t_1, t_2, \dots, t_n) = \prod_{i=1}^n f(t_i)^{\delta_i} S(t_i)^{1-\delta_i} = \prod_{i=1}^n h(t_i)^{\delta_i} S(t_i)$$

C. Loss Function

The following derivation is taken from [3]. For notational convenience, let $\kappa(t) \in \{0, \dots, m\}$ define the index of the discrete time t , meaning $t = \tau_{\kappa(t)}$. Thus, the likelihood contribution for individual i is seen to be

$$\begin{aligned} L_i &= f(t_i)^{\delta_i} S(t_i)^{1-\delta_i} \\ &= [h(t_i) S(\tau_{\kappa(t_i)-1})]^{d_i} [(1 - h(t_i)) S(\tau_{\kappa(t_i)-1})]^{1-d_i} \\ &= h(t_i)^{d_i} [1 - h(t_i)]^{1-d_i} S(\tau_{\kappa(t_i)-1}) \\ &= h(t_i)^{d_i} [1 - h(t_i)]^{1-d_i} \prod_{j=\kappa_{t_i}-1}^{j=1} [1 - h(\tau_j)] \end{aligned}$$

For all the individuals, the combined log-likelihood is

$$L = \prod_{i=1}^n L_i = \prod_{i=1}^n \left(h(t_i)^{d_i} [1 - h(t_i)]^{1-d_i} \prod_{j=\kappa_{t_i}-1}^{j=1} [1 - h(\tau_j)] \right)$$

From here, the loss function for a batch can be constructed as the negative log likelihood.

$$\begin{aligned} \log(L) &= \sum_n^{i=1} \left(d_i \log[h(t_i|x_i)] + (1 - d_i) \log[1 - h(t_i|x_i)] + \right. \\ &\quad \left. \sum_{j=\kappa(t_i)-1}^{\kappa(t_i)-1} j = 1 \log[1 - h(\tau_j|x_i)] \right) \end{aligned}$$

To adjust for varying batch sizes, the mean negative log likelihood is taken as the loss.

$$\begin{aligned} \text{loss} &= -\frac{1}{n} \sum_n^{i=1} \left(d_i \log[h(t_i|x_i)] + (1 - d_i) \log[1 - h(t_i|x_i)] + \right. \\ &\quad \left. \sum_{j=\kappa(t_i)-1}^{j=1} \log[1 - h(\tau_j|x_i)] \right) \\ &= \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^{\kappa_{t_i}} (y_{ij} \log[h(\tau_j|x_i)] + (1 - y_{ij}) \log[1 - h(\tau_j|x_i)]) \end{aligned}$$

This can now be minimized by gradient-based methods, thus making it a useful loss function for a neural

network to work with. Here, y_{ij} is an indicator variable corresponding to 1 if and only if an event is experienced by the individual i at time t_j . Hence, y is a sparse matrix consisting of mostly 0 with 1 only present when the time t_i represents an observed event $d_i = 1$. The loss can be thought of as the negative log likelihood of Bernoulli data $\theta^p (1 - \theta)^{1-p}$, where $\theta = h(\tau_j|x_i)$ and $p = y_{ij}$ (noted by Brown 1975₁) and can be computed using existing functions in the PyTorch library.

D. Neural Network Parameterization

As hazards are conditional probabilities (see Eq. 1), they must lie within $[0, 1]$. A handy function for this is the sigmoid non-linearity

$$g(x) = \frac{1}{1 + e^{-x}}$$

For a neural network taking x_i (the covariates of patient i) as input and producing m outputs denoting m discrete timesteps in the patient's journey, applying the sigmoid function effectively transforms the outputs into valid hazard rates.

$$h(\tau_j|x_i) = g(\phi_j(x_i)) = \frac{1}{1 + e^{-\phi_j(x_i)}}$$

where $\phi_j(x_i)$ is the output of the neural network at node j corresponding to the hazard function over the time period $[m_{j-1}, m_j]$.

4. UNCERTAINTY QUANTIFICATION AND EXPLAINABILITY

An addition this work attempts to make to the above is by introducing a layer of uncertainty quantification via Monte-Carlo dropout [15]. Dropout is generally reserved for training as a means to enforce regularization on the network. When applied during the testing or prediction phase as well, a probability distribution of the output is created, allowing for probabilistic inference instead of point inference as present in regular neural networks. The idea of being uncertainty-aware ties in with the interpretability of modern deep learning methods, which is often a barrier towards the adoption of such methods into highly regulated industries such as finance or medicine. SHAP (SHapley Additive exPlanations) values [16], extended from the field of cooperative game theory provide a unified, model-agnostic method to explain the contributions of covariates to a model's output.

□

-
- [1] G. Savarese and L. H. Lund, Global public health burden of heart failure, *Cardiac Failure Review* **03**, 7 (2017).
 - [2] M. F. Gensheimer and B. Narasimhan, A scalable discrete-time survival model for neural networks, *PeerJ* **7**, 10.7717/peerj.6257 (2019).
 - [3] H. Kvamme and Ø. Borgan, Continuous and discrete-time survival prediction with neural networks (2019), arXiv:1910.06724 [cs, stat].
 - [4] M. Gjoreski, A. Gradisek, B. Budna, M. Gams, and G. Poglajen, Machine learning and end-to-end deep learning for the detection of chronic heart failure from heart sounds, *IEEE Access* **8**, 20313 (2020).
 - [5] J. J. Nirschl, A. Janowczyk, E. G. Peyster, R. Frank, K. B. Margulies, M. D. Feldman, and A. Madabhushi, A deep-learning classifier identifies patients with clinical heart failure using whole-slide images of h&e tissue, *PloS one* **13**, e0192726 (2018).
 - [6] L. Wang, L. Sha, J. R. Lakin, J. Bynum, D. W. Bates, P. Hong, and L. Zhou, Development and Validation of a Deep Learning Algorithm for Mortality Prediction in Selecting Patients With Dementia for Earlier Palliative Care Interventions, *JAMA Network Open* **2**, e196972 (2019).
 - [7] J. R. Ayala Solares, F. E. Diletta Raimondi, Y. Zhu, F. Rahimian, D. Canoy, J. Tran, A. C. Pinho Gomes, A. H. Payberah, M. Zottoli, M. Nazarzadeh, N. Conrad, K. Rahimi, and G. Salimi-Khorshidi, Deep learning for electronic health records: A comparative review of multiple deep neural architectures, *Journal of Biomedical Informatics* **101**, 103337 (2020).
 - [8] G. Lorenzoni, S. S. Sabato, C. Lanera, D. Bottigliengo, C. Minto, H. Ocagli, P. De Paolis, D. Gregori, S. Ilceto, F. Pisanò, and et al., Comparison of machine learning techniques for prediction of hospitalization in heart failure patients, *Journal of Clinical Medicine* **8**, 1298 (2019).
 - [9] Z. Che, Y. Cheng, Z. Sun, and Y. Liu, Exploiting convolutional neural network for risk prediction with medical feature embedding, arXiv preprint arXiv:1701.07474 (2017).
 - [10] Q. Suo, F. Ma, Y. Yuan, M. Huai, W. Zhong, A. Zhang, and J. Gao, Personalized disease prediction using a cnn-based similarity learning method, in *2017 IEEE International Conference on Bioinformatics and Biomedicine (BIBM)* (IEEE, 2017) pp. 811–816.
 - [11] L. Brand, A. Patel, I. Singh, and C. Brand, Real time mortality risk prediction: A convolutional neural network approach., in *HEALTHINF* (2018) pp. 463–470.
 - [12] W. Caicedo-Torres and J. Gutierrez, Iseeu: Visually interpretable deep learning for mortality prediction inside the icu, *Journal of biomedical informatics* **98**, 103269 (2019).
 - [13] C. Nagpal, X. R. Li, and A. Dubrawski, Deep survival machines: Fully parametric survival regression and representation learning for censored data with competing risks (2021), arXiv:2003.01176 [cs, stat].
 - [14] D. F. Moore, Chapter 2 basic principles of survival analysis, in *Applied Survival Analysis Using R* (Springer, 2016).
 - [15] Y. Gal and Z. Ghahramani, Dropout as a bayesian approximation: Representing model uncertainty in deep learning, in *Proceedings of The 33rd International Conference on Machine Learning*, Proceedings of Machine Learning Research, Vol. 48, edited by M. F. Balcan and K. Q. Weinberger (PMLR, New York, New York, USA, 2016) pp. 1050–1059.
 - [16] S. M. Lundberg and S.-I. Lee, A unified approach to interpreting model predictions (2017).